

SEMICON INDIA HACKATHON 2026

Team Alvengers — PS01 Idea Submission

Institution: Vellore Institute of Technology, Vellore

Tiyas Das

Team Leader & ML/AI Lead

Architecture Design &
Loss Refinement

Soumen Mondal

Model Training Engineer

EDSR2x Model
Development &
Training

Partha Protim Mondal

Data Evaluation Engineer

Degradation Forensics
& Metrics Audit

Aryan Raj

Deployment & Presentation Lead

Inference Pipeline &
Packaging

Problem Statement Addressed

AI-Based Restoration of Degraded Images for Semiconductor Inspection_KLA

Semiconductor Inspection & Quality Control Challenge:

- Microscopic inspection images are critical for chip defect detection and yield optimization.
- Low-dose Scanning Electron Microscopy (SEM) suffers from severe coupled image degradations.
- Conventional image filters smooth out critical micro-textures, hiding real wafer defects.

Three Official Degradation Categories (Joint Restoration Challenge):

- 1. Speckle Noise: Random pixel-level noise pushing intensity beyond standard signal bounds
- 2. Gaussian Noise: Reduces edge sharpness, blurring fine line/space grating structures
- 3. Resolution Reduction: Spatial downsampling (128x128 -> 256x256 2x Super-Resolution)

Idea Description: EDSR2x + L1 Loss

Lightweight, dataset-driven single-stage residual deep learning model

Model Architecture Specifications:

- Architecture: EDSR2x Residual CNN (Single-Stage)
- Residual Depth: 8 Residual Blocks (64 Feature Channels)
- Upsampling: 2x PixelShuffle Sub-Pixel Convolution
- Global Skip Connection: bicubic_2x(input) + learned_residual
- Parameters: 776,705 (~0.78M parameters)
- Training Loss: Pure L1 Reconstruction Loss
- Inference Latency: 2.95 ms / image (88.1 FPS on RTX 3050 GPU)

Why EDSR2x Was Selected (Controlled Research):

- ✓ Systematic evaluation across 6 distinct architecture candidates.
- ✓ Proved that complex window attention (SwinIR/HAT) amplified low-SNR noise.
- ✓ Proved that pre-denoising LR images lost high-frequency details.
- ✓ EDSR2x achieved the highest overall PSNR (27.42 dB) and SSIM (0.7357).

Proposed Solution: End-to-End Pipeline


Joint noise removal & 2x super-resolution system architecture

Our Solution: EDSR2x Neural Architecture

Single-stage deep residual network engineered for efficient 2x super-resolution

EDSR2x Model Specifications:

- 8 Residual Blocks (64 Feature Channels)
- 2x PixelShuffle Upsampling Layer
- Global Bicubic Skip Connection
- Parameters: 776,705 (~0.78M params)
- Training Loss: Pure L1 Reconstruction Loss
- Inference Latency: ~2.95 ms / image (RTX 3050 GPU)



Innovation & Uniqueness

Dataset-driven engineering, degradation forensics, and lightweight optimization

1. Read-Only Degradation Forensics

Audited 3,200 paired images. Discovered Poisson-Gaussian shot noise ($\rho = +0.3955$) and direct 2x bicubic downsampling without heavy low-pass optical pre-blur.

2. Controlled Architecture Search

Tested EDSR2x, RCAN-Light, SwinIR-Light, HAT-Small, Wavelet-SR, and Multi-loss. Selected the simplest model with highest measured accuracy.

3. High-Frequency Bottleneck Identification

Identified that high-frequency noise overlap causes single-stage networks to over-smooth dense wafer contact holes, guiding optimal loss selection.

4. Avoiding Over-Processing

Demonstrated that multi-loss (Charbonnier + SSIM + Sobel) caused gradient conflicts (-2.22 dB PSNR loss). Pure L1 reconstruction loss achieved optimal performance.

5. Deployment-Oriented Optimization

0.78M parameters, 111.6 MB peak VRAM, and 2.95 ms GPU latency enable real-time inline semiconductor inspection integration.

Experimental Results & Visual Evidence

Quantitative benchmark and visual restoration comparison

Quantitative Performance Metrics

Rigorous evaluation across 320 official validation & 320 held-out test samples

Evaluation Split	Method	PSNR (dB)	SSIM	MAE	Edge Score	PSNR Win Rate
Validation (320 samples)	Bicubic 2x Baseline	22.59 dB	0.5166	0.0597	0.4727	—
Validation (320 samples)	EDSR2x (Production)	27.42 dB	0.7357	0.0349	0.6639	320/320 (100%)
Validation Net Gain	EDSR2x vs Bicubic	+4.83 dB	+0.2191	-0.0248	+0.1912	100.0%
Held-Out Test (320 samples)	Bicubic 2x Baseline	22.98 dB	0.5243	0.0572	0.4781	—
Held-Out Test (320 samples)	EDSR2x (Production)	27.93 dB	0.7408	0.0310	0.6684	320/320 (100%)
Held-Out Test Net Gain	EDSR2x vs Bicubic	+4.95 dB	+0.2165	-0.0262	+0.1903	100.0%

Technology Stack & Deployment Feasibility

Hardware throughput, standalone inference script, and memory audit

Model Efficiency & Hardware Latency

Engineered for real-time semiconductor line inspection integration



Designed for efficient, reproducible, high-throughput wafer defect inspection

GitHub Repository & Demo Video

Open-source codebase, reproducible packaging, and presentation video

Official Public GitHub Repository:

https://github.com/TiyasDas-81/Semicon_Hackathon_26.git

Official Presentation Demo Video:

File Name: EDSR2x_demo.mp4 (Included in fresh_training/final_submission/)

Video Attributes: 83 seconds duration | 1920x1080 Full HD | 30 FPS H.264 MP4

Evaluation Script Usage Command:

`python evaluation.py --input <test_images_directory> --output <restored_output_directory>`

References & Academic Citations

Authoritative literature and official competition problem specifications

1. Lim, B., Son, S., Kim, H., Nah, S., & Lee, K. M. (2017). Enhanced deep residual networks for single image super-resolution. In IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW), pp. 136-144.
2. Zhang, Y., Li, K., Li, K., Wang, L., Zhong, B., & Fu, Y. (2018). Image super-resolution using very deep residual channel attention networks. In Proceedings of the European Conference on Computer Vision (ECCV), pp. 286-301.
3. Liang, J., Cao, J., Sun, G., Zhang, K., Van Gool, L., & Timofte, R. (2021). SwinIR: Image restoration using swin transformer. In IEEE/CVF International Conference on Computer Vision Workshops (ICCVW), pp. 1833-1844.
4. Chen, X., Wang, X., Zhou, J., & Dong, C. (2023). Activating more pixels in image super-resolution transformer. In IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pp. 22367-22377.
5. Paszke, A., et al. (2019). PyTorch: An imperative style, high-performance deep learning library. In Advances in Neural Information Processing Systems (NeurIPS 32), pp. 8026-8037.
6. i4C & KLA SemiCon India Hackathon 2026. Problem Statement PS01: AI-Based Restoration of Degraded Images for Semiconductor Inspection.